

# UNSUPERVISED HIERARCHICAL CONCEPT LEARNING

**Sumegh Roychowdhury\***

Indian Institute of Technology, Kharagpur

**Sumedh A. Sontakke\***

University of Southern California

**Nikaash Puri, Mausoom Sarkar, Milan Aggarwal, Pinkesh Badjatiya, Balaji Krishnamurthy**  
Adobe Media and Data Science Research

**Laurent Itti**

University of Southern California

## ABSTRACT

Discovering concepts (or temporal abstractions) in an unsupervised manner from demonstration data is an important problem. Organizing these discovered concepts hierarchically at different levels of abstraction is useful in discovering patterns, building ontologies, and generating tutorials from demonstration data. However, it is very challenging to discover such concepts and relationships (or a hierarchy) between these discovered concepts, purely from demonstration data, without any access to the environment. In this paper, we present a Transformer-based concept abstraction architecture UNHCLE (pronounced uncle) that extracts a hierarchy of concepts in an unsupervised way from demonstration data. We empirically demonstrate how UNHCLE discovers meaningful hierarchies using datasets from Chess and Cooking domains. Finally, we show how UNHCLE learns meaningful language labels for concepts by using demonstration data augmented with natural language for cooking and chess. All of our code is available at <https://github.com/UNHCLE/UNHCLE/>.

## 1 INTRODUCTION


Consider a video (Figure 1) that demonstrates how to cook an egg. Humans subconsciously learn concepts (such as boiling water) that describe different concepts (or skills) in such demonstrations Pammi et al. (2004). These learned skills can be composed and reused in different ways to learn new concepts. Discovering such concepts automatically from demonstration data is a non-trivial problem. Shankar et al. (2019) introduces a sequence-to-sequence architecture that clusters long-horizon action trajectories into shorter temporal skills. However, their approach treats skills as independent concepts. In contrast, humans organize these concepts in hierarchies where lower-level concepts can be grouped to define higher-level concepts Naim et al. (2019).

We extend the architecture in Shankar et al. (2019) to simultaneously discover concepts along with their hierarchical organization without any supervision.

We propose an end-to-end trainable architecture UNHCLE for hierarchical representation learning from demonstrations. UNHCLE takes as input a long horizon trajectory of high-dimensional images demonstrating a complex task (in our case, chess and cooking) and the associated textual commentary and isolates semantically meaningful subsequences in input trajectories. We emphasize that it does not require temporal annotations which link subsequences in the trajectories of images to the free-flowing commentary, but instead, autonomously discovers this mapping. Therefore, this work takes a step towards unsupervised video understanding of high-dimensional data. Our contributions can be summarized as follows:

- We introduce a transformer-based architecture to learn a multi-modal hierarchical latent embedding space to encode the various concepts in long-horizon demonstration trajectories.

\*Equal contributions. Work done as research interns at Adobe Media and Data Science Research Team.



**Figure 1:** Overview of our approach. UNHCLE learns a hierarchical latent space of concepts describing long horizon tasks like cooking and chess gameplay.

UNHCLE abstracts these concepts (shown through visual qualitative analysis) without requiring any temporal supervision, i.e., it divides long-horizon trajectories into semantically meaningful subsequences, without access to any temporal annotations that split these trajectories optimally.

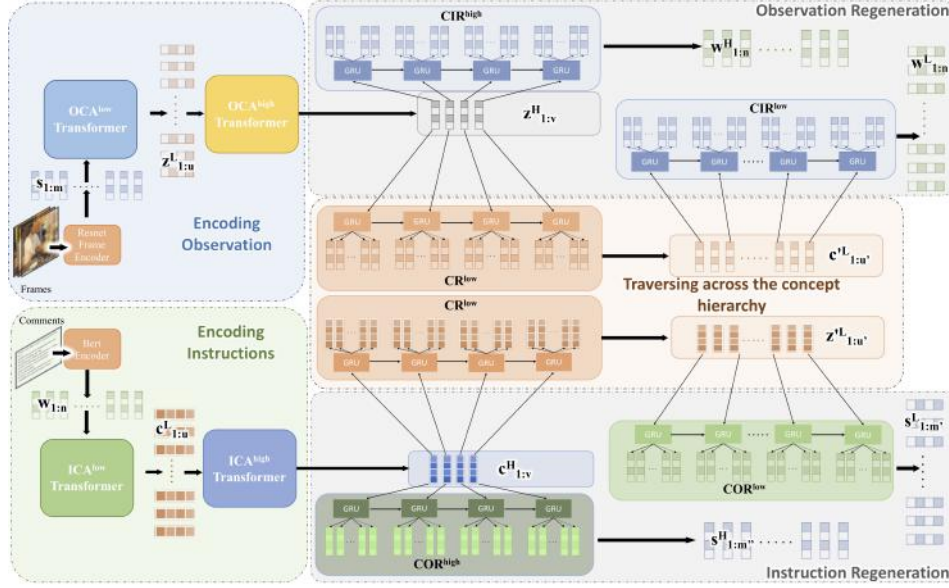
- We show the quantitative effectiveness of learning high-level concepts in a hierarchical manner compared to learning them in isolation while outperforming several baselines on YouCook2 (Zhou et al. (2017)) and Chess Opening dataset<sup>1</sup>.
- We further introduce a mechanism to incorporate commentary accompanying demonstrations in UNHCLE and show improvements in hierarchical concepts discovered.
- We introduce TimeWarped IoU (TW-IoU), an evaluation metric that we use to compare the alignment of our discovered concepts and ground-truth events.

Existing approaches to representation learning for demonstrations or videos typically require significant supervision. Typically, sequence-to-sequence architectures are trained on datasets segmented by humans. During inference, these architectures generate proposals for timestamps that segment the input trajectory into semantically meaningful sequences. These complex sequence-to-sequence models require significant amounts of annotated data, making them costly to train them.

More generally, video and scene understanding is an important research area with wide-ranging applications. Most recently, Chen et al. (2019) utilize semantic awareness to perform complex depth estimation tasks to acquire the geometric properties of 3-dimensional space from 2-dimensional images. Tosi et al. (2020) utilize similar semantic information for depth estimation, optical flow and motion segmentation. Boggust et al. (2019) attempt to ground words in the video, but apply significant supervision to synchronize them, requiring human intervention. We attempt to learn similar embeddings but do so in a completely unsupervised manner, not utilizing any of the temporal labels available.

The field of learning from demonstrations (Nicolescu & Mataric (2003)) seeks to learn to perform tasks from a set of demonstrated behaviors. Behavioral Cloning is one popular scheme (Esmaili et al. (1995)). Atkeson & Schaal (1997) and Schaal (1997) show how agents can learn simple tasks like cartpole simply from demonstrations. Pastor et al. (2009) also study how robots can learn from human demonstrations of tasks. Peters et al. (2013) and Kober & Peters (2009) fit a parametric model to the demonstrations. Niekum et al. (2012), Murali et al. (2016), and Meier et al. (2011) first segment trajectories into subsequence and then apply a parametric models to each subsequence. More recently, Schmeckpeper et al. (2019) shows that agents can learn to maximize external reward using a large corpus of observation data, i.e., trajectories of states on a relatively smaller corpus of interaction data, i.e., trajectories of state-action pairs. Hierarchical task representations have been studied as well. Instead of treating demonstrations in a flat manner, one may also infer the hierarchical structure. A few recent works attempt to do so (Xu et al. (2018); Sun et al. (2018)), or as task graphs (Huang et al., 2019). Both Xu et al. (2018) and Huang et al. (2019) address generalizing to new instances of manipulation tasks in the few-shot regime by abstracting away low-level controls. However, all of these approaches require an environment i.e., a transition and reward function to learn from. On the contrary, humans show an ability to learn by watching demonstrations, which we attempt to replicate.

<sup>1</sup><https://www.kaggle.com/residentmario/recommending-chess-openings>



**Figure 2:** Overview of our approach. UNHCLE learns a hierarchical semantically-meaningful embedding space which allows it to perform complex downstream tasks such as temporal concept segmentation and label prediction. Refer to Section 2.1 for details about the abbreviations used in the figure.

Temporal abstractions of action sequences or skill/primitive learning is also a related field. Eysenbach et al. (2018), learn a large number of low-level sequences of actions by forcing the agent to produce skills that are different from those previously acquired. However, due to the diversity bias, the agent results in learning many useless skills that cannot be used for any semantically meaningful task. Similarly, Sharma et al. (2019) attempt to learn skills such that their transitions are almost deterministic in a given environment. These approaches also require access to an environment, whereas we try to learn without an environment.

## 2 APPROACH

### 2.1 UNHCLE: UNSUPERVISED HIERARCHICAL CONCEPT LEARNING

Intuitively, we define a concept as a short sequence of states which repeatedly occur across several demonstration trajectories. Concepts have an upper limit on their length in time-steps. These concepts can be obtained from the images of demonstrations, denoted by  $z$ , and from the associated textual description, represented by  $c$ . We also refer to them as Observational Concept (OC) and Instructional Concept (IC) respectively and the module that encodes these concepts is referred to as Observational Concept Abstraction (OCA) and Instructional Concept Abstraction (ICA) modules. Additionally, concepts are hierarchical in nature - thus, lower-level and higher-level observational or image concepts are denoted by  $z^L$  and  $z^H$ , respectively. Analogously, lower-level, and higher-level textual concepts are represented by  $c^L$  and  $c^H$ , respectively. Similarly, the higher-level and lower-level modules are denoted by *low* or *high* for the corresponding level.

Once we obtain these concept abstractions across levels, we can also transform them across the levels using a Concept Regeneration (CR) Module to transform *low*-level concepts to *high*-level concepts and vice-versa. Instead of just traversing in the concept level hierarchy, we can also use the Concept Instruction Regeneration Module or CIR to obtain the original instructions that map to that concept in its respective concept modality.

Subsequently, we provide details about the different stages in our proposed technique.

**Encoding Observation:** Given a long horizon trajectory of demonstration images along with its associated textual description, UNHCLE is able to abstract a hierarchy of concepts from demonstration images. We first pass these input images through ResNet-32 (He et al. (2016)) to get a sequence of

image vectors as  $S = s_{1:m}$ , and the associated text is converted into word vectors  $W = w_{1:n}$  using BERT-base (Devlin et al. (2018)). Observations combine to produce lower-level concepts whereas higher-level concepts are simply aggregations of such lower-level concepts. Thus, the Lower-level Observation Concept Abstraction module ( $OCA^{low}$ ) is trained to embed a sequence of image vectors of a video ( $s_1, s_2, \dots, s_m$ ) into a sequence of concept vectors ( $z_1^L, z_2^L, \dots, z_U^L$ ) where  $u \ll m$  such that  $z_{1:u}^L = OCA^{low}(s_{1:m})$ . Subsequently, lower-level concepts combine together to form a higher level of concepts using the Higher-level Observation Concept Abstraction module ( $OCA^{high}$ ) such that  $z_{1:v}^H = OCA^{high}(z_{1:u}^L)$ .

**Encoding Instructions:** We also endeavour to discover these higher-level and lower-level concepts through natural language instructions. The Lower-level Instruction Concept Abstraction module ( $ICA^{low}$ ) and Higher-level Instruction Concept Abstraction module ( $ICA^{high}$ ) are responsible for this functionality. From a corpus of words, ( $w_1, w_2, \dots, w_n$ ), the  $ICA^{low}$  module generates concept ( $c_1^L, c_2^L, \dots, c_u^L$ ),  $u \ll n$ :  $c_{1:u}^L = ICA^{low}(w_{1:n})$ . Subsequently the  $ICA^{high}$  encodes the lower-level language concepts  $c_{1:u}^L$  into higher level concepts as  $c_{1:v}^H = ICA^{high}(c_{1:u}^L)$ .

**Traversing across the concept hierarchy:** Learning concepts at different levels has an added advantage of traversing these concepts in the concept hierarchy. This additionally allows us to utilize these hierarchical traversals to obtain coarse or fine-grained concepts at any level. We can thus regenerate the lower concepts from higher-level instruction concepts using the Lower-level Concept Regeneration Module ( $CR^{low}$ ) such that  $z_{1:u'}^L = CR^{low}(c_{1:v}^H)$ . We can then later utilize this to obtain lower-level concepts from the higher-level concepts to regenerate the demonstration images  $S^L = s_{1:m}^L$  in a cross-modal fashion.

**Observation and Instruction Regeneration:** Under a concept, the sequence of frames is nearly deterministic i.e. the knowledge of a concept uniquely identifies the accompanying sequence of images in the demonstration trajectory. Subsequently, we regenerate the demonstration image vectors  $S^L = s_{1:m}^L$  from lower-level concepts using Lower-level Concept Observation Regeneration Module ( $COR^{low}$ ) such that  $s_{1:m}^L = COR^{low}(z_{1:u'}^L)$ . We also regenerate the demonstration image vectors  $S^U = s_{1:m}^U$  from higher-level concepts abstracted from language using Higher-level Concept Observation Regeneration Module ( $COR^{high}$ ) such that  $s_{1:m}^U = COR^{high}(c_{1:v}^H)$  in a similar cross-modal manner.

Finally, inspired by humans who can easily describe concept representations using free-flowing natural language, we first regenerate lower-level concepts from higher-level observation concepts using the Lower-level Concept Regeneration Module ( $CR^{low}$ ) such that  $c_{1:u'}^L = CR^{low}(z_{1:v}^H)$ , and subsequently, regenerate the word vectors  $W^L = w_{1:n}^L$  from lower-level concepts using Lower-level Concept Instruction Regeneration Module ( $CIR^{low}$ ) such that  $w_{1:n}^L = CIR^{low}(c_{1:u'}^L)$ . Additionally, the higher-level concepts identified by the  $OCA^{high}$  module from demonstration frames are also described using a meaningful free-flowing commentary by the Higher-level Concept Instruction Regeneration module or  $CIR^{high}$ . Thus, we regenerate the word vectors  $W^U = w_{1:n}^U$  from higher-level concepts using  $CIR^{high}$  such that  $w_{1:n}^U = CIR^{high}(z_{1:v}^H)$ .

## 2.2 SOFT-DYNAMIC TIME WARPING (SOFT-DTW)

The most significant challenge that our work overcomes is that we do not require supervision for hierarchical temporal segmentation i.e. we do not require annotations which demarcate the beginning and ending of a concept, both in language and in the space of frame's timestamps. Currently, most approaches rely on annotation describing the frame start and end for an event and the associated text for this event. Our architecture is trained using free-flowing commentary without event demarcations. This is achieved using the loss function we implement. We utilize Soft-DTW (Cuturi & Blondel (2017)) loss to align the reconstructed trajectories with the input long-horizon trajectory. So given two trajectories  $x = (x_1, x_2, \dots, x_n)$  and  $y = (y_1, y_2, \dots, y_m)$ , the soft-DTW( $x, y$ ) is computed as,

$$\text{soft-DTW}(x, y) = \min^\gamma \{ \langle \mathcal{A}, \Delta(x, y) \rangle, \mathcal{A} \} \quad (1)$$

where  $\mathcal{A} \in \mathcal{A}_{n,m}$  is the alignment matrix,  $\Delta(x, y) = [\delta(x_i, y_i)]_{ij} \in \mathbb{R}^{n \times m}$  and  $\delta$  being the cost function.  $\min^\gamma$  operator is then computed as,

$$\min^\gamma \{a_1, \dots, a_n\} = \begin{cases} \min_{i \leq n} a_i, & \gamma = 0, \\ -\gamma \log \sum_{i=1}^n e^{-a_i/\gamma}, & \gamma > 0. \end{cases} \quad (2)$$

For our experiments, we use  $L_2$  distance as  $\delta$  and  $\gamma = 1$ .

### 2.3 LEARNING OBJECTIVE

Our approach that has been outlined in Figure 2 uses several loss terms between different network outputs to achieve our objective. The soft-DTW is used between several sequences as follows.

$$\mathcal{L}_{dyn} = \text{soft-DTW}(z_{1:u}^L, c_{1:u'}^L) + \text{soft-DTW}(c_{1:u}^L, z_{1:u'}^L) + \text{soft-DTW}(s_{1:m}^L, s_{1:m'}^L) + \text{soft-DTW}(s_{1:m}^H, s_{1:m''}^H)$$

We use the Negative Log-Likelihood (NLL) loss ( $\mathcal{L}_{nll}$ ) between the generated comment vectors and the BERT vectors and use the  $L_2$  loss between concepts from different modalities as follows,

$$\mathcal{L}_{static} = \mathcal{L}_{nll}(w_{1:n}, w_{1:n}^L) + \mathcal{L}_{nll}(w_{1:n}, w_{1:n}^H) + L_2(z_{1:v}^H, c_{1:v}^H) + L_2(z_{1:u'}^L, c_{1:u'}^L)$$

We then define our final loss as,  $\mathcal{L}_{total} = \alpha * \mathcal{L}_{dyn} + \beta * \mathcal{L}_{static}$

### 2.4 EVALUATION METRICS

The ground-truth events in the dataset and the concepts generated by UNHCLE may differ in number, duration, and start-time. To evaluate the efficacy of UNHCLE in generating concepts that align with the human-annotated events in our dataset, it is imperative that we utilize a metric that measures the overlap between generated concepts and ground truths and also accounts for this possible temporal mismatch. To this end, we introduce time-warp IoU. Note that we do not utilize these the temporal ground-truth event annotations during training, but only to compare the abstraction generated by our architecture to the human-annotated events.

Consider the search series  $S = (s_1, s_2, s_3 \dots s_M) \in \mathbb{S}$  and target series  $T = (t_1, t_2, t_3 \dots t_N)$  where  $S$  corresponds to the end-of-concept time stamp for each concept as generated by UNHCLE for a single long-horizon demonstration trajectory. Thus, the  $i^{th}$  concept abstracted from UNHCLE starts at time =  $s_{i-1}$  and end at time =  $s_i$ . Similarly,  $T$  corresponds to the end-of-event time stamp for each ground-truth event in the demonstration trajectory, where the  $j^{th}$  ground truth event starts at time =  $t_{j-1}$  and ends at time =  $t_j$ . Note that both  $s_0$  and  $t_0$  are equal to zero i.e. we measure time starting at zero for all demonstration trajectories.

#### 2.4.1 TIME-WARPED IOU

We calculate the time-warped alignment between the event annotations and the concepts from UNHCLE. This implies calculating  $\Delta(S, T)$ , solving the following optimization problem (Berndt & Clifford (1994)),  $\Delta(S, T) = \min_{P \in \mathcal{P}} \sum_{m,n \in P} \delta(s_m, t_n)$ , where the  $S$  and  $T$  correspond to the search and target series respectively and  $\delta$  corresponds to a distance metric (in our case the  $L_2$  norm), measuring time mismatch.

$\Delta(S, T)$  therefore corresponds to the trajectory discrepancy measure defined as the matching cost for the optimal matching path  $P$  among all possible valid matching paths  $\mathcal{P}$  (i.e. paths satisfying monotonicity, continuity, and boundary conditions). From this optimal trajectory obtained we can also get the warping function  $W$  such that  $W(s_i) = t_j$ , i.e. we find the optimal mapping between the  $i^{th}$  concept ending at time =  $s_i$  and the  $j^{th}$  event ending at time =  $t_j$ . To quantify the quality of this mapping we introduce Time-Warp Intersection over Union for a single long-horizon trajectory  $S$ .

$$TW - IoU(S) = \sum_{t_i} \frac{\sum_{s_j: W(s_j)=t_i} \min(t_i, s_j) - \max(t_{i-1}, s_{j-1})}{\max_{s_j: W(s_j)=t_i} \{\max(t_i, s_j)\} - \min_{s_j: W(s_j)=t_i} \{\min(t_{i-1}, s_{j-1})\}} \quad (3)$$

Intuitively, this corresponds to the overlap between the  $i^{th}$  concept ending at time =  $s_i$  and the  $j^{th}$  event ending at time =  $t_j$ .

### 3 EXPERIMENTS

#### 3.1 DATASETS

In order to find conceptual temporal patterns from videos in a hierarchical manner, we experiment with two datasets from different domains.

**YouCook2** (Zhou et al. (2017)) dataset comprises of instructional videos for 89 unique recipes (~22 videos per recipe) containing labels that separate the long horizon trajectories of demonstrations into events - with explicit time stamps for the beginning and end of each event along with the associated commentary. It contains 1,333 videos for training and 457 videos for testing. The average number of segments per video is 7.7 and the average duration of the video is 5.27 minutes.

**Recommending Chess Openings**<sup>2</sup> dataset consists of opening moves in the game of Chess. An Opening in Chess is a fixed sequence of moves which when performed leads to a final board state putting the player in a strategic position in the game. Commonly used chess openings are each labeled with a name (Figure 3 shows examples). The dataset contains 20,058 openings with each containing a sequence of chess moves and it’s corresponding opening and variation labels. The train-test split used for our experiments is 80-20.

#### 3.2 IMPLEMENTATION DETAILS

The Observational Concept Abstraction (OCA) and Instructional Concept Abstraction (ICA) modules consist of the Transformer (Vaswani et al. (2017)) Encoder with 8 hidden layers and 8-Head Attention which takes as input, a sequence of frames positionally encoded and outputs a hidden vector. It is then passed through a Transformer Decoder again having 8 hidden layers and encoder-decoder cross-attention with masking which finally outputs the concepts which are latent variables having dimension  $conceptlength \times 768$ . We use a 1-layer GRU (Chung et al. (2014)) for our Concept Instruction Regeneration (CIR), Concept Observation Regeneration (COR), and the Concept Regeneration (CR) module which converts the high-level concepts to low-level concepts. For encoding the image frames, we down-sample all videos to 500 frames and use ResNet-32 (pretrained on MSCOCO dataset) (He et al. (2016)) for dense embedding of dimension  $512 \times 1$ . We further down-sample to 200 frames per trajectory for all experiments. Comments are encoded using BERT-base pre-trained embeddings with a 768 hidden dimension.

For our low-level concepts, we keep the maximum number of concepts discovered as 16 and for high-level concepts as 4. These assumptions are based on the YouCook2 dataset statistics wherein ground-truth annotations (not used during training) the minimum number of segments were 5 and the maximum as 16. We keep  $\alpha = 1$  and  $\beta = 1$  for all our experiments along with a batch-size of 128 and use Adam optimizer for training with a learning rate of  $1e - 5$  and train the network for 100 epochs.


#### 3.3 RESULTS AND ANALYSIS

##### 3.3.1 VISUALIZING HIERARCHY

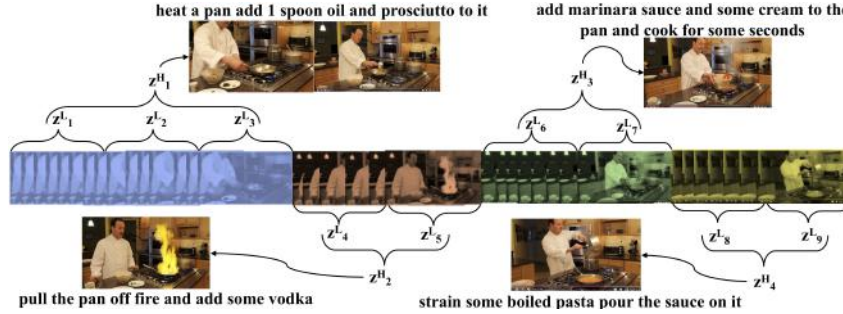
Here we analyze whether the discovered concepts are human interpretable i.e. *are the temporal clusters within a single demonstration semantically meaningful?*. We find that our architecture abstracts several useful human interpretable concepts without any supervision. Figure 3 and 4 shows the results obtained. These abstracted high-level concepts align well with the ground truth event labels provided in YouCook2. We also additionally find that our model is able to split each ground truth event into lower-level concepts. For instance, in a pasta-making demonstration in YouCook2, a single event corresponding to the description “*heat a pan add 1 spoon oil and prosciutto to it*”, UNHCLE is divided into low level concepts corresponding to “*heat pan*”, “*add oil*” and “*prosciutto*”. Note that no explicit event time labels were provided to UNHCLE, indicating that our model is able to abstract such coherent sub-sequences, thus taking the first step towards video understanding.

We also visualize the concepts abstracted by UNHCLE when trained on chess opening data. The concepts learnt here also produce coherent, human-interpretable results.

<sup>2</sup><https://www.kaggle.com/residentmario/recommending-chess-openings>



**Figure 3:** Hierarchy of concepts discovered by UNHCLE using openings data in Chess. At a high level, UNHCLE correctly identifies concepts corresponding to the Blackmar-Diemer Gambit and the Pietrowsky Defense. At a low level, it identifies concepts such as “*d4 d5..* and *e4 d3 ..*” that are used across several openings



**Figure 4:** Example Hierarchy of concepts discovered by UNHCLE on the YouCook2 dataset

### 3.3.2 COMPARISON WITH BASELINES

In this section, we evaluate the performance of UNHCLE quantitatively on YouCook2<sup>3</sup> and quantify its ability to generate coherent concepts that align with the human annotated ground truths using the TW-IoU metric. We compare our approach with 4 baselines. (1) **Random** baseline predicts segment randomly on the basis of uniformly sampled timestamps for a given video on the basis of its duration. (2) **EQUALDIV** consists of dividing the video into conceptual segments of equal duration (3) **GRU-supervised**: Further we also consider a naive supervised baseline comprising of a GRU Cho et al. (2014) based encoder that sequentially processes the Resnet features corresponding to frames in a video followed by a decoder GRU that attends Bahdanau et al. (2014) on encoder outputs and is trained to sequentially predict end timestamp of each meaningful segment (variable in number) in the video. (4) **FLAT**: We implement the Shankar et al. (2019) approach which takes as input, a sequence of video frames and discovers a single level of concepts without any hierarchy.

For baselines, we predict only one level of concepts and hence refer to them as high-level concepts. Table 1a summarises and compares the TW-IoU computed between ground truth time stamp annotations and predicted/discovered segments. As can be seen, UNHCLE achieves the best TW-IoU compared to all other baselines. It can be noted that discovering only one (high) level of concepts (FLAT) results in low TW-IoU (for High-Level Concepts). However, UNHCLE discovers concepts that align better with the ground truth events (UNHCLE performs  $\sim 23\%$  better) compared to FLAT. Further, the GRU-supervised baseline performs better than FLAT but UNHCLE outperforms ( $\sim 15\%$ ) the GRU-supervised baseline which shows that predicting such concepts is difficult even for a supervised baseline.

### 3.3.3 EFFECT OF GUIDANCE THROUGH COMMENTARY

In this section, we discuss the improvement in the quality and precision of conceptual segments discovered by UNHCLE using video commentary to guide our model training on YouCook2 dataset. To study this, we compare UNHCLE without comment which discovers concept hierarchy using just frame observations in a video against UNHCLE with comment which additionally uses commentary as guide (as in Figure 2). Table 1a compares the TW-IoU of segments obtained corresponding to

<sup>3</sup>We use YouCook2 since each recipe comprises of multiple conceptual segments unlike Chess dataset (each game is just an opening + variation) which we instead use for evaluating label prediction ability later.



Method	TW-IoU	
Random	28.2	
EQUALDIV	31.7	
GRU-supervised	22.8	
FLAT	14.4	
	<i>high-level</i>	<i>low-level</i>
UNHCLE w/o commentary	37.4	58.7
UNHCLE w/ commentary	<b>47.4</b>	<b>59.7</b>


(a)

Frames Used	TW-IoU	
	<i>high</i>	<i>low</i>
200	37.4	58.7
64	33.4	34.0
32	13.1	12.7

(b)

**Table 1:** (a) TW-IoU scores for single level concepts predicted by the baselines along with the TW-IoU scores for both the high-level and low-level concepts for predictions from our proposed technique UNHCLE (b) Comparison showing the trade-off between number of down-sampled frames and TW-IoU for UNHCLE

High Level Concepts for each of these two models. As can be seen, using commentary significantly improves the TW-IoU by  $\sim 10\%$  which establishes that using commentary enables UNHCLE to detect precise boundaries for segments corresponding to various concepts. However, we find only a marginal improvement of  $\sim 1\%$  in TW-IoU corresponding to low level concepts. Further, it can be seen in Figure 5 that using commentary allows the model to better understand how to merge lower-level concepts appropriately into higher-level concepts - “*heat a pan, add 1 spoon oil and prosciutto to it*” separately without merging it with subsequent segments.



**Figure 5:** In this figure, we show that using commentary as guide during training, UNHCLE learns to better combine low-level concepts to form high-level concepts which are better aligned towards our ground-truth annotations. A refers to the concept segments discovered w/o comments and B refers to the one with comment.

### 3.3.4 LABEL PREDICTION TASK FOR CHESS OPENINGS

We further evaluate UNHCLE on the task of label (name of opening + variation) prediction corresponding to hierarchical concepts discovered. Since recommending Chess Openings comprises of a label for each sequence of moves in an opening, we train UNHCLE w/ commentary (Figure 2). Since there are at max only two segments corresponding to opening move and variation move in the dataset we use **label prediction accuracy** as our metric here instead of *TW-IoU*. There are 300 distinct labels and our model achieves **78.2%** accuracy which shows that it is able to correctly discover the correct segmentation between openings with different variations as hierarchical concepts and associate them with the correct label in most of the test samples.

### 3.3.5 EFFECT OF SAMPLING RATE ON THE QUALITY OF HIERARCHY

For YouCook2 we down-sample 200 frames from the original 500 frames provided in the dataset and use that as our input to UNHCLE. We analyze the trade-off between time and performance if we further down-sample the frames. This also provides us a better insight into how much granular information regarding the video do we need to be able to discover better hierarchies. Table 1b shows



the results obtained for our experiment. Interestingly, we don't observe a linear drop in performance by reducing the number of frames. Even with just 64 frames, the results beat the other baselines.

## 4 CONCLUSION

We show that our approach (UNHCLE) discovers concepts and organizes them in a meaningful hierarchy using only demonstration data from chess openings and cooking. We also show that this discovered hierarchy of concepts is useful in predicting textual labels and temporal concept segmentations for the associated demonstrations. It would be interesting to use the concept hierarchy discovered by UNHCLE to generate a curriculum where lower-level concepts would be taught first followed by higher-level concepts. Another good direction would be to extend UNHCLE to learn different types of relationships between concepts without supervision.

## REFERENCES

- Christopher G Atkeson and Stefan Schaal. Robot learning from demonstration. In *ICML*, volume 97, pp. 12–20. Citeseer, 1997.
- Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*, 2014.
- Donald J Berndt and James Clifford. Using dynamic time warping to find patterns in time series. In *KDD workshop*, volume 10, pp. 359–370. Seattle, WA, USA:, 1994.
- Angie W Boggust, Kartik Audhkhasi, Dhiraj Joshi, David Harwath, Samuel Thomas, Rogério Schmidt Feris, Danny Gutfreund, Yang Zhang, Antonio Torralba, Michael Picheny, et al. Grounding spoken words in unlabeled video. In *CVPR Workshops*, pp. 29–32, 2019.
- Po-Yi Chen, Alexander H Liu, Yen-Cheng Liu, and Yu-Chiang Frank Wang. Towards scene understanding: Unsupervised monocular depth estimation with semantic-aware representation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 2624–2632, 2019.
- Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using RNN encoder–decoder for statistical machine translation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 1724–1734, Doha, Qatar, October 2014. Association for Computational Linguistics. doi: 10.3115/v1/D14-1179. URL <https://www.aclweb.org/anthology/D14-1179>.
- Junyoung Chung, Caglar Gulcehre, KyungHyun Cho, and Yoshua Bengio. Empirical evaluation of gated recurrent neural networks on sequence modeling. *arXiv preprint arXiv:1412.3555*, 2014.
- Marco Cuturi and Mathieu Blondel. Soft-DTW: a differentiable loss function for time-series. volume 70 of *Proceedings of Machine Learning Research*, pp. 894–903, International Convention Centre, Sydney, Australia, 06–11 Aug 2017. PMLR.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.
- Nasser Esmaili, Claude Sammut, and GM Shirazi. Behavioural cloning in control of a dynamic system. In *1995 IEEE International Conference on Systems, Man and Cybernetics. Intelligent Systems for the 21st Century*, volume 3, pp. 2904–2909. IEEE, 1995.
- Benjamin Eysenbach, Abhishek Gupta, Julian Ibarz, and Sergey Levine. Diversity is all you need: Learning skills without a reward function. *arXiv preprint arXiv:1802.06070*, 2018.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778, 2016.

- De-An Huang, Suraj Nair, Danfei Xu, Yuke Zhu, Animesh Garg, Li Fei-Fei, Silvio Savarese, and Juan Carlos Niebles. Neural task graphs: Generalizing to unseen tasks from a single video demonstration. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 8565–8574, 2019.
- Jens Kober and Jan Peters. Learning motor primitives for robotics. In *2009 IEEE International Conference on Robotics and Automation*, pp. 2112–2118. IEEE, 2009.
- Franziska Meier, Evangelos Theodorou, Freek Stulp, and Stefan Schaal. Movement segmentation using a primitive library. In *2011 IEEE/RSJ International Conference on Intelligent Robots and Systems*, pp. 3407–3412. IEEE, 2011.
- Adithyavairavan Murali, Animesh Garg, Sanjay Krishnan, Florian T Pokorny, Pieter Abbeel, Trevor Darrell, and Ken Goldberg. Tsc-dl: Unsupervised trajectory segmentation of multi-modal surgical demonstrations with deep learning. In *2016 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 4150–4157. IEEE, 2016.
- Michelangelo Naim, Mikhail Katkov, Stefano Recanatesi, and Misha Tsodyks. Emergence of hierarchical organization in memory for random material. *Scientific Reports*, 9(1):1–10, 2019.
- Monica N Nicolescu and Maja J Mataric. Natural methods for robot task learning: Instructive demonstrations, generalization and practice. In *Proceedings of the second international joint conference on Autonomous agents and multiagent systems*, pp. 241–248, 2003.
- Scott Niekum, Sarah Osentoski, George Konidaris, and Andrew G Barto. Learning and generalization of complex tasks from unstructured demonstrations. In *2012 IEEE/RSJ International Conference on Intelligent Robots and Systems*, pp. 5239–5246. IEEE, 2012.
- VS Chandrasekhar Pammi, Krishna P Miyapuram, Raju S Bapi, and Kenji Doya. Chunking phenomenon in complex sequential skill learning in humans. In *International Conference on Neural Information Processing*, pp. 294–299. Springer, 2004.
- Peter Pastor, Heiko Hoffmann, Tamim Asfour, and Stefan Schaal. Learning and generalization of motor skills by learning from demonstration. In *2009 IEEE International Conference on Robotics and Automation*, pp. 763–768. IEEE, 2009.
- Jan Peters, Jens Kober, Katharina Mülling, Oliver Krämer, and Gerhard Neumann. Towards robot skill learning: From simple skills to table tennis. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, pp. 627–631. Springer, 2013.
- Stefan Schaal. Learning from demonstration. In *Advances in neural information processing systems*, pp. 1040–1046, 1997.
- Karl Schmeckpeper, Annie Xie, Oleh Rybkin, Stephen Tian, Kostas Daniilidis, Sergey Levine, and Chelsea Finn. Learning predictive models from observation and interaction. *arXiv preprint arXiv:1912.12773*, 2019.
- Tanmay Shankar, Shubham Tulsiani, Lerrel Pinto, and Abhinav Gupta. Discovering motor programs by recomposing demonstrations. In *International Conference on Learning Representations*, 2019.
- Archit Sharma, Shixiang Gu, Sergey Levine, Vikash Kumar, and Karol Hausman. Dynamics-aware unsupervised discovery of skills. *arXiv preprint arXiv:1907.01657*, 2019.
- Shao-Hua Sun, Hyeonwoo Noh, Sriram Somasundaram, and Joseph Lim. Neural program synthesis from diverse demonstration videos. In *International Conference on Machine Learning*, pp. 4790–4799, 2018.
- Fabio Tosi, Filippo Aleotti, Pierluigi Zama Ramirez, Matteo Poggi, Samuele Salti, Luigi Di Stefano, and Stefano Mattoccia. Distilled semantics for comprehensive scene understanding from videos. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4654–4665, 2020.

- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in neural information processing systems*, pp. 5998–6008, 2017.
- Danfei Xu, Suraj Nair, Yuke Zhu, Julian Gao, Animesh Garg, Li Fei-Fei, and Silvio Savarese. Neural task programming: Learning to generalize across hierarchical tasks. In *2018 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 1–8. IEEE, 2018.
- Luowei Zhou, Chenliang Xu, and Jason J Corso. Towards automatic learning of procedures from web instructional videos. *arXiv preprint arXiv:1703.09788*, 2017.